

Global Air Quality Analysis - IDS Semester Project

Reg# 2023-CS-646

```
In [4]: # Import Libraries and Create Output Folders
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from scipy.stats import zscore
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             confusion_matrix, classification_report,
                             mean_absolute_error, mean_squared_error, r2_score)
import os
import warnings
warnings.filterwarnings('ignore')

# Create folders for saving graphs
folders = ['graphs/1_preprocessing', 'graphs/2_eda_univariate', 'graphs/3_eda_bivariate',
           'graphs/4_correlation', 'graphs/5_comparative', 'graphs/6_timeseries']
for folder in folders:
    os.makedirs(folder, exist_ok=True)

plt.style.use('seaborn-v0_8-whitegrid')
print("✓ Libraries imported and folders created!")

✓ Libraries imported and folders created!
```

```
In [5]: # Load Dataset
df = pd.read_csv('global_air_quality_dataset.csv')
df_original = df.copy()

print(f"Shape: {df.shape}")
print(f"\nColumns: {list(df.columns)}")
print(f"\nFirst 5 rows:")
df.head()
```

Shape: (10000, 12)

Columns: ['City', 'Country', 'Date', 'PM2.5', 'PM10', 'NO2', 'SO2', 'CO', 'O3', 'Temperature', 'Humidity', 'Wind Speed']

First 5 rows:

Out[5]:

	City	Country	Date	PM2.5	PM10	NO2	SO2	CO	O3	Temperature	H
0	Bangkok	Thailand	2023-03-19	86.57	25.19	99.88	30.63	4.46	36.29		17.67
1	Istanbul	Turkey	2023-02-16	50.63	97.39	48.14	8.71	3.40	144.16		3.46
2	Rio de Janeiro	Brazil	2023-11-13	130.21	57.22	98.51	9.92	0.12	179.31		25.29
3	Mumbai	India	2023-03-16	119.70	130.52	10.96	33.03	7.74	38.65		23.15
4	Paris	France	2023-04-04	55.20	36.62	76.85	21.85	2.00	67.09		16.02



In [6]:

```
# Dataset Info and Statistics
print("Data Types:")
print(df.dtypes)
print("\nStatistical Summary:")
df.describe()
```

Data Types:

City	object
Country	object
Date	object
PM2.5	float64
PM10	float64
NO2	float64
SO2	float64
CO	float64
O3	float64
Temperature	float64
Humidity	float64
Wind Speed	float64
dtype:	object

Statistical Summary:

Out[6]:

	PM2.5	PM10	NO2	SO2	CO
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	77.448439	104.438161	52.198649	25.344490	5.047984
std	41.927871	55.062396	27.320490	14.091194	2.852625
min	5.020000	10.000000	5.010000	1.000000	0.100000
25%	41.185000	57.137500	28.347500	13.190000	2.560000
50%	77.725000	103.690000	52.100000	25.350000	5.090000
75%	113.392500	152.265000	75.705000	37.500000	7.480000
max	149.980000	200.000000	100.000000	49.990000	10.000000



1. Data Preprocessing & Cleaning

```
In [7]: # 1.1 Check and Handle Missing Values
print("Missing Values:")
print(df.isnull().sum())

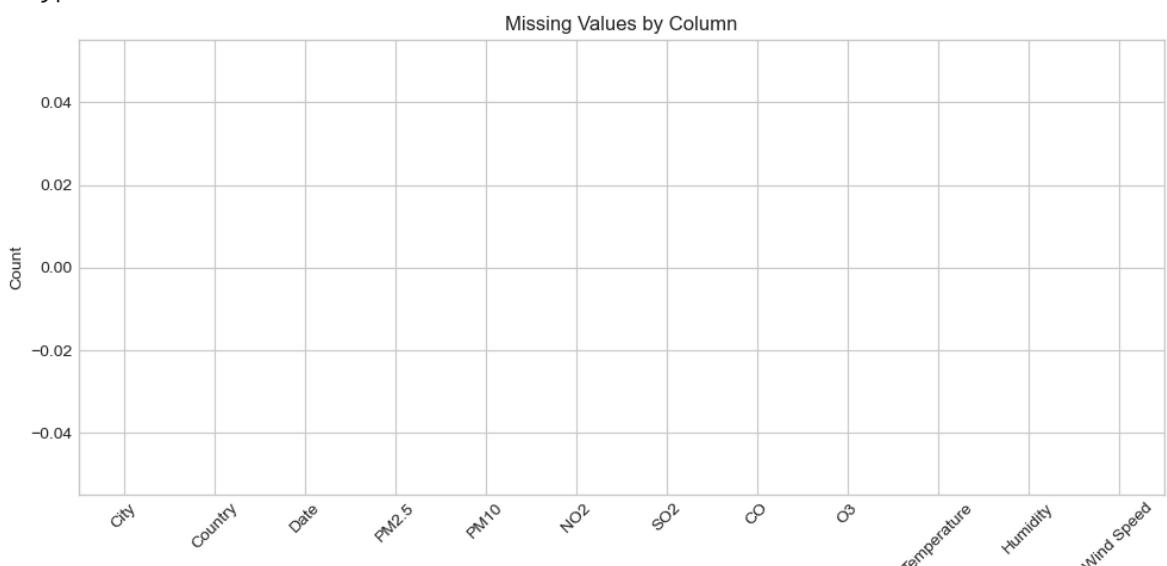
# Visualize missing values
fig, ax = plt.subplots(figsize=(10, 5))
df.isnull().sum().plot(kind='bar', color=['red' if x > 0 else 'green' for x in df.isnull().sum()])
plt.title('Missing Values by Column')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig('graphs/1_preprocessing/missing_values.png', dpi=150)
plt.show()

# Impute missing values
numerical_cols = ['PM2.5', 'PM10', 'NO2', 'SO2', 'CO', 'O3', 'Temperature', 'Humidity', 'Wind Speed']
for col in numerical_cols:
    df[col].fillna(df[col].median(), inplace=True)

print("\n✓ Missing values after imputation:", df.isnull().sum().sum())
```

Missing Values:

```
City          0
Country       0
Date          0
PM2.5         0
PM10          0
NO2           0
SO2           0
CO            0
O3            0
Temperature   0
Humidity      0
Wind Speed    0
dtype: int64
```



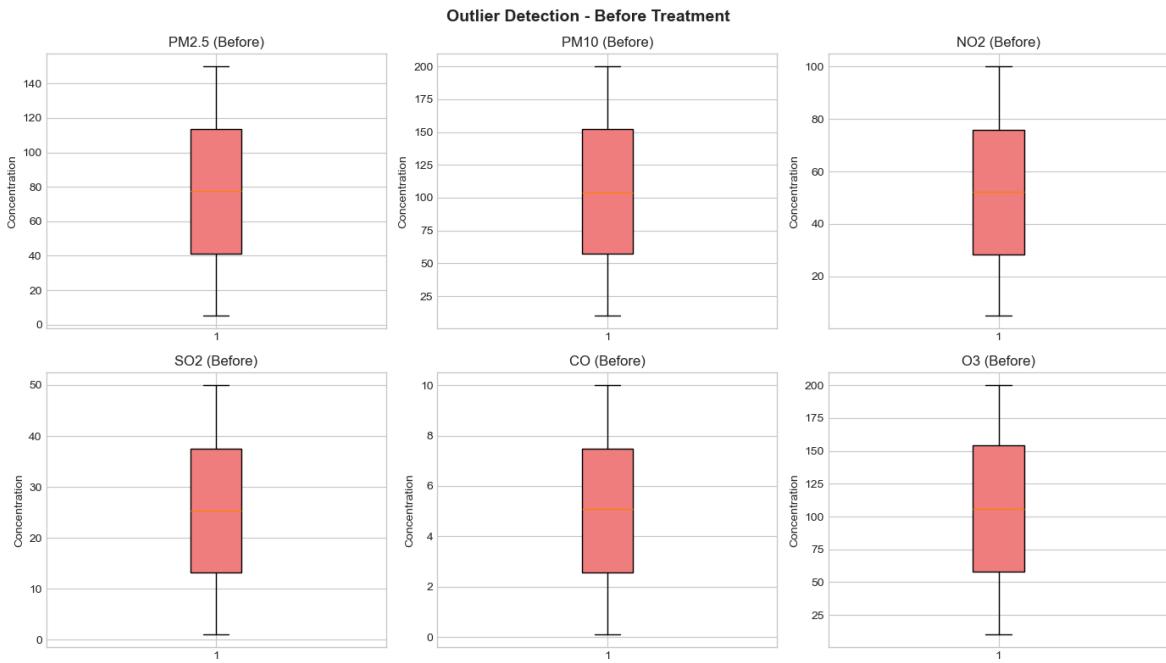
✓ Missing values after imputation: 0

```
In [8]: # 1.2 Outlier Detection - Box Plots Before Treatment
pollutant_cols = ['PM2.5', 'PM10', 'NO2', 'SO2', 'CO', 'O3']
```

```

fig, axes = plt.subplots(2, 3, figsize=(14, 8))
axes = axes.flatten()
for i, col in enumerate(pollutant_cols):
    axes[i].boxplot(df[col], patch_artist=True, boxprops=dict(facecolor='lightcoral'))
    axes[i].set_title(f'{col} (Before)')
    axes[i].set_ylabel('Concentration')
plt.suptitle('Outlier Detection - Before Treatment', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/1_preprocessing/outliers_before.png', dpi=150)
plt.show()

```



```

In [9]: # 1.3 Outlier Detection using IQR and Z-Score Methods
print("Outlier Detection Summary:")
print("-" * 60)

for col in numerical_cols:
    # IQR Method
    Q1, Q3 = df[col].quantile(0.25), df[col].quantile(0.75)
    IQR = Q3 - Q1
    iqr_outliers = ((df[col] < Q1 - 1.5*IQR) | (df[col] > Q3 + 1.5*IQR)).sum()

    # Z-Score Method
    z_outliers = (np.abs(zscore(df[col])) > 3).sum()

    print(f"{col}: IQR Outliers={iqr_outliers}, Z-Score Outliers={z_outliers}")

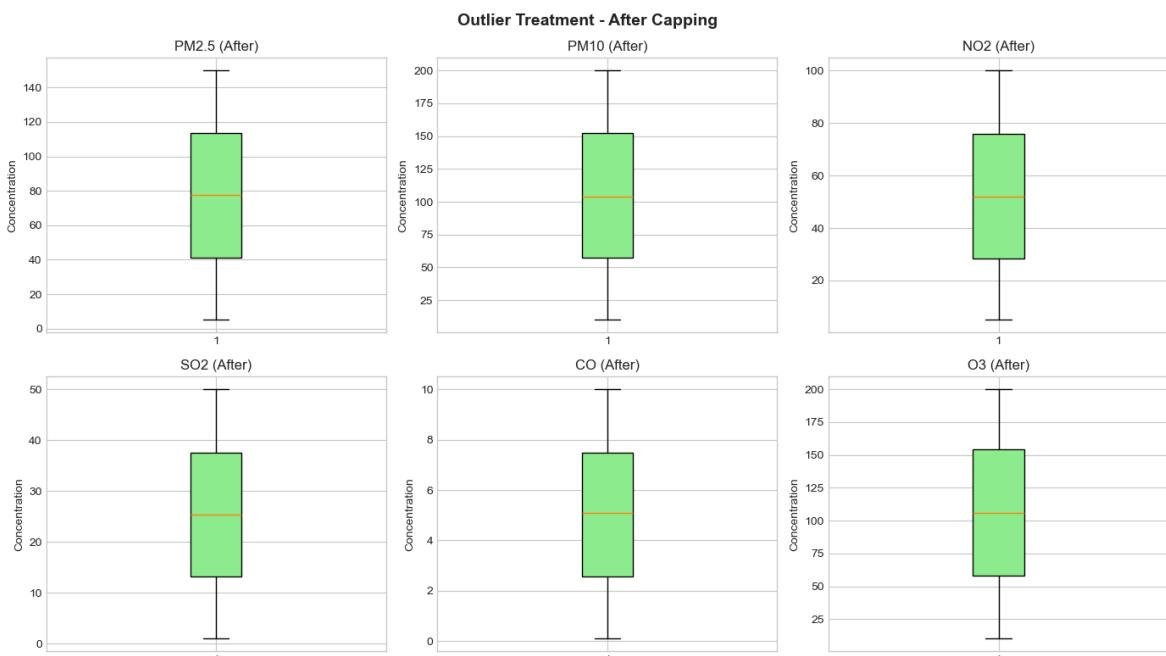
```

Outlier Detection Summary:

PM2.5: IQR Outliers=0, Z-Score Outliers=0
 PM10: IQR Outliers=0, Z-Score Outliers=0
 NO2: IQR Outliers=0, Z-Score Outliers=0
 SO2: IQR Outliers=0, Z-Score Outliers=0
 CO: IQR Outliers=0, Z-Score Outliers=0
 O3: IQR Outliers=0, Z-Score Outliers=0
 Temperature: IQR Outliers=0, Z-Score Outliers=0
 Humidity: IQR Outliers=0, Z-Score Outliers=0
 Wind Speed: IQR Outliers=0, Z-Score Outliers=0

```
In [10]: # 1.4 Handle Outliers using IQR Capping
for col in numerical_cols:
    Q1, Q3 = df[col].quantile(0.25), df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower, upper = Q1 - 1.5*IQR, Q3 + 1.5*IQR
    df[col] = df[col].clip(lower, upper)

# Box Plots After Treatment
fig, axes = plt.subplots(2, 3, figsize=(14, 8))
axes = axes.flatten()
for i, col in enumerate(pollutant_cols):
    axes[i].boxplot(df[col], patch_artist=True, boxprops=dict(facecolor='lightgray'))
    axes[i].set_title(f'{col} (After)')
    axes[i].set_ylabel('Concentration')
plt.suptitle('Outlier Treatment - After Capping', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/1_preprocessing/outliers_after.png', dpi=150)
plt.show()
print("✓ Outliers handled using IQR capping")
```



- ✓ Outliers handled using IQR capping

```
In [11]: # 1.5 Calculate AQI and Create Categories
pm25_bp = [(0, 12, 0, 50), (12.1, 35.4, 51, 100), (35.5, 55.4, 101, 150),
            (55.5, 150.4, 151, 200), (150.5, 250.4, 201, 300), (250.5, 500, 301,
            500)

def calc_aqi(pm25):
    for low, high, aqi_low, aqi_high in pm25_bp:
        if low <= pm25 <= high:
            return ((aqi_high - aqi_low) / (high - low)) * (pm25 - low) + aqi_low
    return 500

df['AQI'] = df['PM2.5'].apply(calc_aqi)

# Categorize AQI
def aqi_category(aqi):
    if aqi <= 50: return 'Good'
    elif aqi <= 100: return 'Moderate'
    elif aqi <= 150: return 'Unhealthy'
    else: return 'Hazardous'
```

```

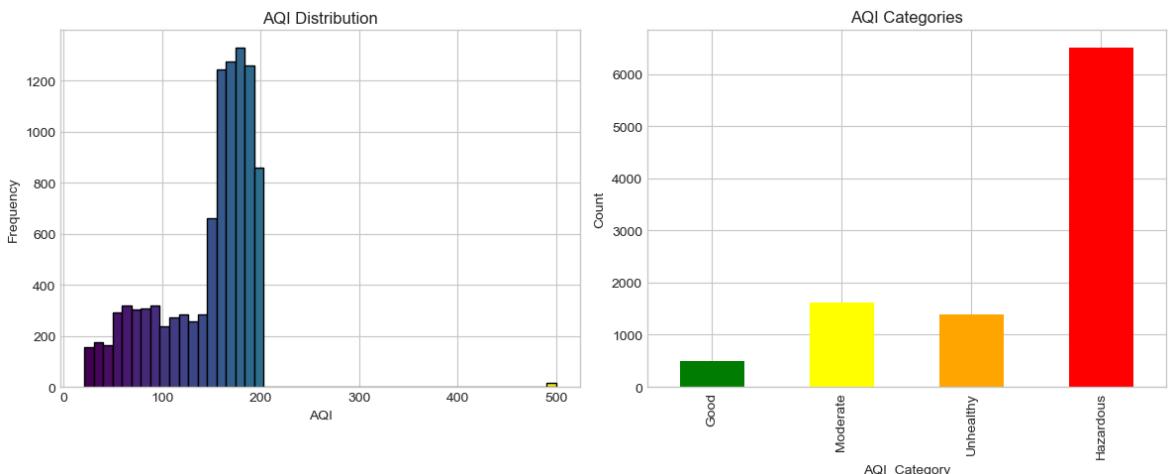
df['AQI_Category'] = df['AQI'].apply(aqi_category)
df['AQI_Encoded'] = df['AQI_Category'].map({'Good': 0, 'Moderate': 1, 'Unhealthy': 2, 'Hazardous': 3})

# Visualize AQI Distribution
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
n, bins, patches = axes[0].hist(df['AQI'], bins=50, edgecolor='black')
cmap = plt.cm.viridis
for i, p in enumerate(patches):
    p.set_facecolor(cmap(i / len(patches)))
axes[0].set_title('AQI Distribution')
axes[0].set_xlabel('AQI')
axes[0].set_ylabel('Frequency')

colors = ['green', 'yellow', 'orange', 'red']
df['AQI_Category'].value_counts().reindex(['Good', 'Moderate', 'Unhealthy', 'Hazardous'], kind='bar', ax=axes[1], color=colors)
axes[1].set_title('AQI Categories')
axes[1].set_ylabel('Count')
plt.tight_layout()
plt.savefig('graphs/1_preprocessing/aqi_distribution.png', dpi=150)
plt.show()

print("AQI Category Distribution:")
print(df['AQI_Category'].value_counts())

```



AQI Category Distribution:

AQI_Category	Count
Hazardous	6514
Moderate	1609
Unhealthy	1380
Good	497

Name: count, dtype: int64

```

In [12]: # 1.6 Feature Scaling - Standardization and Normalization
features = ['PM2.5', 'PM10', 'NO2', 'SO2', 'CO', 'O3', 'Temperature', 'Humidity']

# Compare scaling methods
scaler_std = StandardScaler()
scaler_mm = MinMaxScaler()

df_std = df.copy()
df_norm = df.copy()
df_std[features] = scaler_std.fit_transform(df[features])
df_norm[features] = scaler_mm.fit_transform(df[features])

```

```

# Enhanced Visualization
fig, axes = plt.subplots(2, 3, figsize=(16, 10))

# Row 1: PM2.5 Scaling Comparison
axes[0,0].hist(df['PM2.5'], bins=40, color="#e74c3c", alpha=0.7, edgecolor='black')
axes[0,0].axvline(df['PM2.5'].mean(), color='darkred', linestyle='--', linewidth=2)
axes[0,0].set_title('Original PM2.5', fontsize=12, fontweight='bold')
axes[0,0].set_xlabel('Concentration ( $\mu\text{g}/\text{m}^3$ )')
axes[0,0].set_ylabel('Frequency')
axes[0,0].legend()

axes[0,1].hist(df_std['PM2.5'], bins=40, color="#3498db", alpha=0.7, edgecolor='black')
axes[0,1].axvline(0, color='darkblue', linestyle='--', linewidth=2, label='Mean')
axes[0,1].set_title('Standardized (Z-score)', fontsize=12, fontweight='bold')
axes[0,1].set_xlabel('Standard Deviations from Mean')
axes[0,1].set_ylabel('Frequency')
axes[0,1].legend()

axes[0,2].hist(df_norm['PM2.5'], bins=40, color="#2ecc71", alpha=0.7, edgecolor='black')
axes[0,2].axvline(df_norm['PM2.5'].mean(), color='darkgreen', linestyle='--', linewidth=2)
axes[0,2].set_title('Normalized (Min-Max)', fontsize=12, fontweight='bold')
axes[0,2].set_xlabel('Scaled Value (0-1)')
axes[0,2].set_ylabel('Frequency')
axes[0,2].legend()

# Row 2: Before/After comparison for multiple features
sample_features = ['PM2.5', 'Temperature', 'Humidity']
colors = ['#e74c3c', '#f39c12', '#9b59b6']

# Original scale boxplot
bp1 = axes[1,0].boxplot([df[f] for f in sample_features], patch_artist=True, labels=sample_features)
for patch, color in zip(bp1['boxes'], colors):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)
axes[1,0].set_title('Original Scale', fontsize=12, fontweight='bold')
axes[1,0].set_ylabel('Value (different units)')

# Standardized boxplot
bp2 = axes[1,1].boxplot([df_std[f] for f in sample_features], patch_artist=True, labels=sample_features)
for patch, color in zip(bp2['boxes'], colors):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)
axes[1,1].axhline(0, color='black', linestyle='--', alpha=0.5)
axes[1,1].set_title('After Standardization', fontsize=12, fontweight='bold')
axes[1,1].set_ylabel('Z-score ( $\sigma$  units)')

# Normalized boxplot
bp3 = axes[1,2].boxplot([df_norm[f] for f in sample_features], patch_artist=True, labels=sample_features)
for patch, color in zip(bp3['boxes'], colors):
    patch.set_facecolor(color)
    patch.set_alpha(0.7)
axes[1,2].set_title('After Normalization', fontsize=12, fontweight='bold')
axes[1,2].set_ylabel('Scaled Value (0-1)')

plt.suptitle('Feature Scaling Comparison: Standardization vs Normalization', fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/1_preprocessing/scaling_comparison.png', dpi=150, bbox_inches='tight')
plt.show()

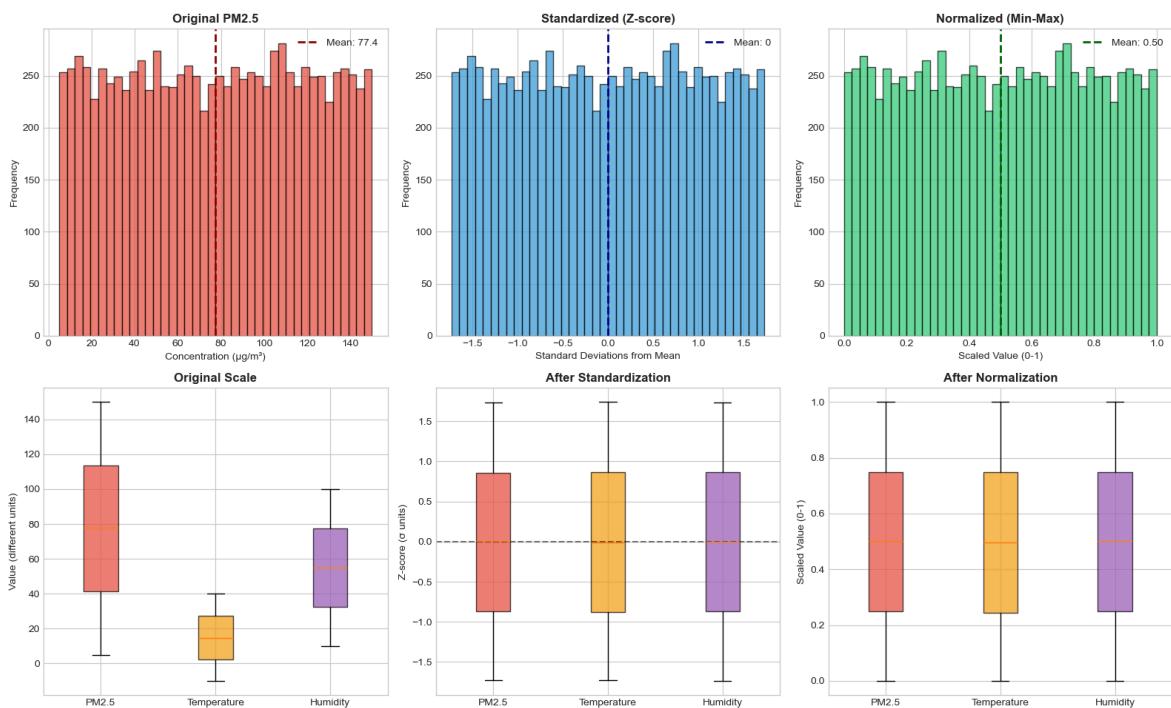
```

```

# Print summary statistics
print("📊 Scaling Methods Summary:")
print("-" * 60)
print(f"{'Method':<20} {'Mean':<15} {'Std':<15} {'Range':>15}")
print("-" * 60)
print(f"{'Original PM2.5':<20} {df['PM2.5'].mean():<15.2f} {df['PM2.5'].std():<15.2f} {[5.0, 150.0]}")
print(f"{'Standardized':<20} {df_std['PM2.5'].mean():<15.2f} {df_std['PM2.5'].std():<15.2f} {[ -1.73, 1.73]}")
print(f"{'Normalized':<20} {df_norm['PM2.5'].mean():<15.2f} {df_norm['PM2.5'].std():<15.2f} {[0.00, 1.00]}")
print("-" * 60)
print("✓ StandardScaler: Mean=0, Std=1 (useful for algorithms sensitive to scale)")
print("✓ MinMaxScaler: Range [0,1] (preserves original distribution shape)")

```

Feature Scaling Comparison: Standardization vs Normalization



📊 Scaling Methods Summary:

Method	Mean	Std	Range
<hr/>			
Original PM2.5	77.45	41.93	[5.0, 150.0]
Standardized	-0.00	1.00	[-1.73, 1.73]
Normalized	0.50	0.29	[0.00, 1.00]

- ✓ StandardScaler: Mean=0, Std=1 (useful for algorithms sensitive to scale)
- ✓ MinMaxScaler: Range [0,1] (preserves original distribution shape)

In [13]: # 1.7 Train-Test Split (80-20)

```

X = df[features]
y = df['AQI_Encoded']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print(f"Training set: {len(X_train)} samples ({len(X_train)/len(X)*100:.0f}%)")
print(f"Test set: {len(X_test)} samples ({len(X_test)/len(X)*100:.0f}%)")

```

```
print(f"\nClass distribution in training set:")
print(pd.Series(y_train).value_counts().sort_index())
```

Training set: 8000 samples (80%)
Test set: 2000 samples (20%)

Class distribution in training set:

AQI_Encoded

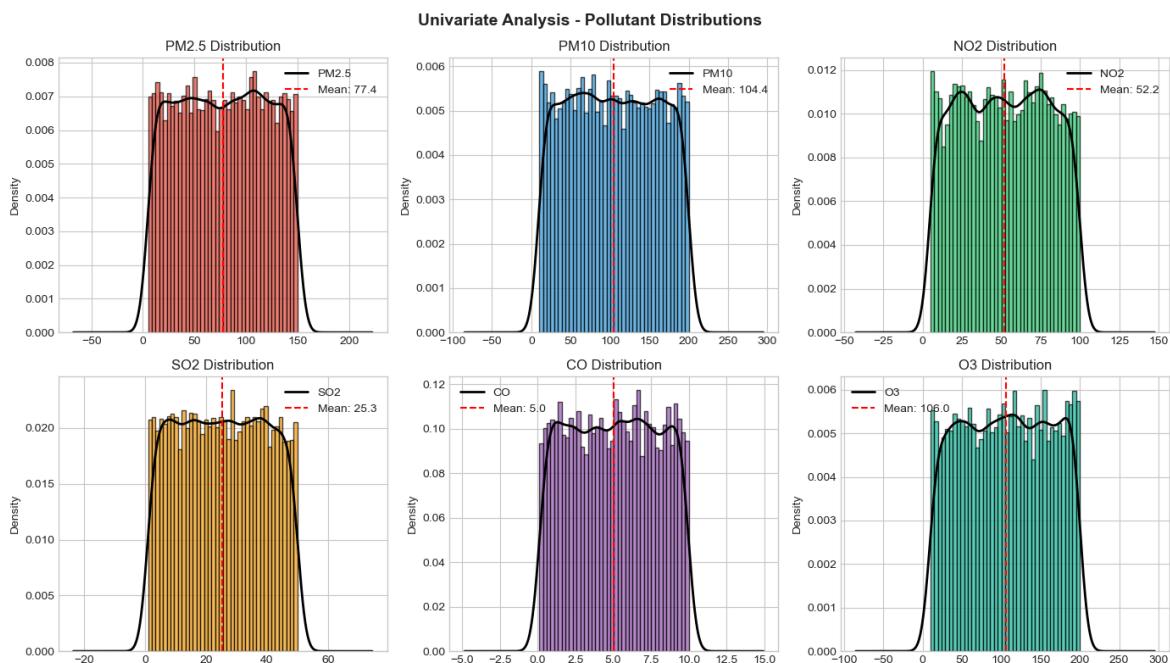
0	398
1	1287
2	1104
3	5211

Name: count, dtype: int64

2. Exploratory Data Analysis (EDA)

```
In [14]: # 2.1 Univariate Analysis - Pollutant Distributions
fig, axes = plt.subplots(2, 3, figsize=(14, 8))
axes = axes.flatten()
colors = ['#e74c3c', '#3498db', '#2ecc71', '#f39c12', '#9b59b6', '#1abc9c']

for i, (col, color) in enumerate(zip(pollutant_cols, colors)):
    axes[i].hist(df[col], bins=40, color=color, alpha=0.7, edgecolor='black', density=True)
    df[col].plot(kind='kde', ax=axes[i], color='black', linewidth=2)
    axes[i].axvline(df[col].mean(), color='red', linestyle='--', label=f'Mean: {df[col].mean():.2f}')
    axes[i].set_title(f'{col} Distribution')
    axes[i].legend()
plt.suptitle('Univariate Analysis - Pollutant Distributions', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/2_eda_univariate/pollutant_distributions.png', dpi=150)
plt.show()
```



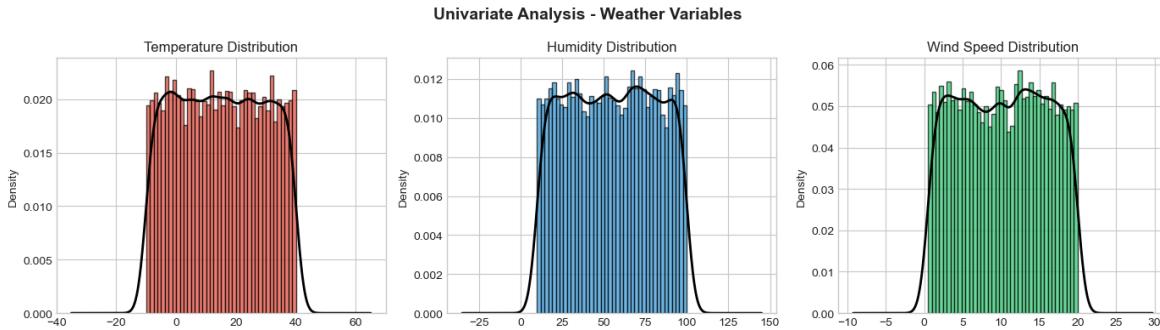
```
In [15]: # 2.1b Univariate Analysis - Weather Variables
fig, axes = plt.subplots(1, 3, figsize=(14, 4))
weather_cols = ['Temperature', 'Humidity', 'Wind Speed']
weather_colors = ['#e74c3c', '#3498db', '#2ecc71']

for i, (col, color) in enumerate(zip(weather_cols, weather_colors)):
```

```

        axes[i].hist(df[col], bins=40, color=color, alpha=0.7, edgecolor='black', density=True)
        df[col].plot(kind='kde', ax=axes[i], color='black', linewidth=2)
        axes[i].set_title(f'{col} Distribution')
    plt.suptitle('Univariate Analysis - Weather Variables', fontsize=14, fontweight='bold')
    plt.tight_layout()
    plt.savefig('graphs/2_eda_univariate/weather_distributions.png', dpi=150)
    plt.show()

```



```

In [16]: # 2.1c Skewness and Kurtosis Analysis
print("Skewness and Kurtosis Analysis:")
print("-" * 50)
for col in numerical_cols:
    skew = df[col].skew()
    kurt = df[col].kurtosis()
    print(f"{col}: Skewness={skew:.3f}, Kurtosis={kurt:.3f}")

```

Skewness and Kurtosis Analysis:

```

PM2.5: Skewness=-0.006, Kurtosis=-1.204
PM10: Skewness=0.011, Kurtosis=-1.201
NO2: Skewness=-0.002, Kurtosis=-1.192
SO2: Skewness=0.007, Kurtosis=-1.192
CO: Skewness=-0.002, Kurtosis=-1.199
O3: Skewness=-0.014, Kurtosis=-1.196
Temperature: Skewness=0.013, Kurtosis=-1.201
Humidity: Skewness=-0.006, Kurtosis=-1.201
Wind Speed: Skewness=-0.008, Kurtosis=-1.209

```

```

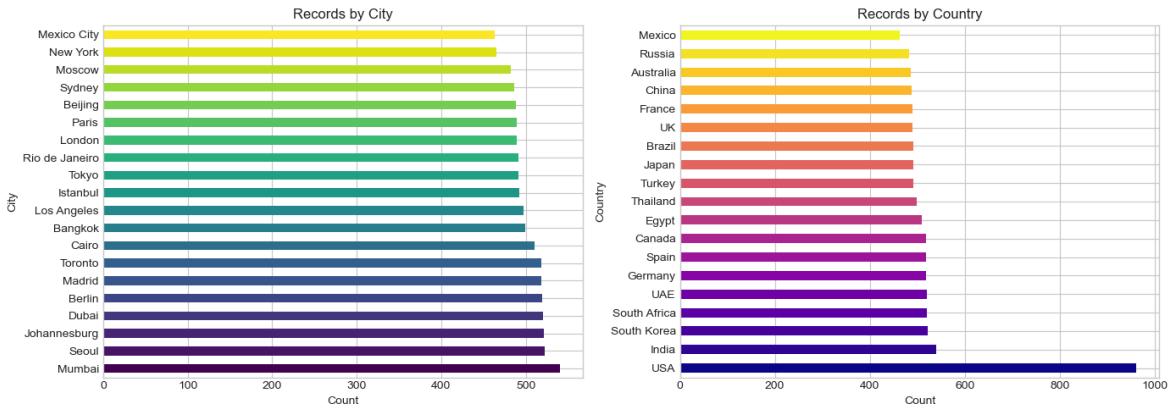
In [17]: # 2.1d Categorical Variables Distribution
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

df['City'].value_counts().plot(kind='barh', ax=axes[0], color=plt.cm.viridis(np.linspace(0, 1, len(df['City'].unique()))))
axes[0].set_title('Records by City')
axes[0].set_xlabel('Count')

df['Country'].value_counts().plot(kind='barh', ax=axes[1], color=plt.cm.plasma(np.linspace(0, 1, len(df['Country'].unique()))))
axes[1].set_title('Records by Country')
axes[1].set_xlabel('Count')

plt.tight_layout()
plt.savefig('graphs/2_eda_univariate/categorical_distribution.png', dpi=150)
plt.show()

```



```
In [18]: # 2.2 Bivariate Analysis - Pollutants vs Weather (Line Charts with Trend)
fig, axes = plt.subplots(2, 3, figsize=(16, 10))

# Helper function to create binned line chart with trend
def plot_bivariate_line(ax, x_col, y_col, color, xlabel, ylabel, title):
    # Create bins for x variable
    df_temp = df[[x_col, y_col]].dropna()
    df_temp['bin'] = pd.cut(df_temp[x_col], bins=20)
    grouped = df_temp.groupby('bin')[y_col].agg(['mean', 'std']).reset_index()
    grouped['bin_mid'] = grouped['bin'].apply(lambda x: x.mid)
    grouped = grouped.dropna()

    # Plot Line with confidence band
    ax.plot(grouped['bin_mid'], grouped['mean'], marker='o', linewidth=2, color=color)
    ax.fill_between(grouped['bin_mid'],
                    grouped['mean'] - grouped['std']/2,
                    grouped['mean'] + grouped['std']/2,
                    alpha=0.2, color=color)

    # Add trend line
    z = np.polyfit(grouped['bin_mid'], grouped['mean'], 1)
    p = np.poly1d(z)
    ax.plot(grouped['bin_mid'], p(grouped['bin_mid']), '--', color='black', linewidth=2)

    # Calculate correlation
    corr = df[x_col].corr(df[y_col])
    ax.set_xlabel(xlabel, fontsize=10)
    ax.set_ylabel(ylabel, fontsize=10)
    ax.set_title(f'{title}\n(r = {corr:.3f})', fontsize=11, fontweight='bold')
    ax.legend(loc='best', fontsize=8)
    ax.grid(True, alpha=0.3)

# Row 1: PM2.5 relationships
plot_bivariate_line(axes[0,0], 'Temperature', 'PM2.5', '#e74c3c',
                    'Temperature (°C)', 'PM2.5 (µg/m³)', 'PM2.5 vs Temperature')
plot_bivariate_line(axes[0,1], 'Humidity', 'PM2.5', '#3498db',
                    'Humidity (%)', 'PM2.5 (µg/m³)', 'PM2.5 vs Humidity')
plot_bivariate_line(axes[0,2], 'Wind Speed', 'PM2.5', '#2ecc71',
                    'Wind Speed (m/s)', 'PM2.5 (µg/m³)', 'PM2.5 vs Wind Speed')

# Row 2: Other pollutants
plot_bivariate_line(axes[1,0], 'Temperature', 'O3', '#f39c12',
                    'Temperature (°C)', 'O3 (ppb)', 'O3 vs Temperature')
plot_bivariate_line(axes[1,1], 'Temperature', 'NO2', '#9b59b6',
                    'Temperature (°C)', 'NO2 (ppb)', 'NO2 vs Temperature')
plot_bivariate_line(axes[1,2], 'Wind Speed', 'CO', '#1abc9c',
                    'Wind Speed (m/s)', 'CO (ppm)', 'CO vs Wind Speed')
```

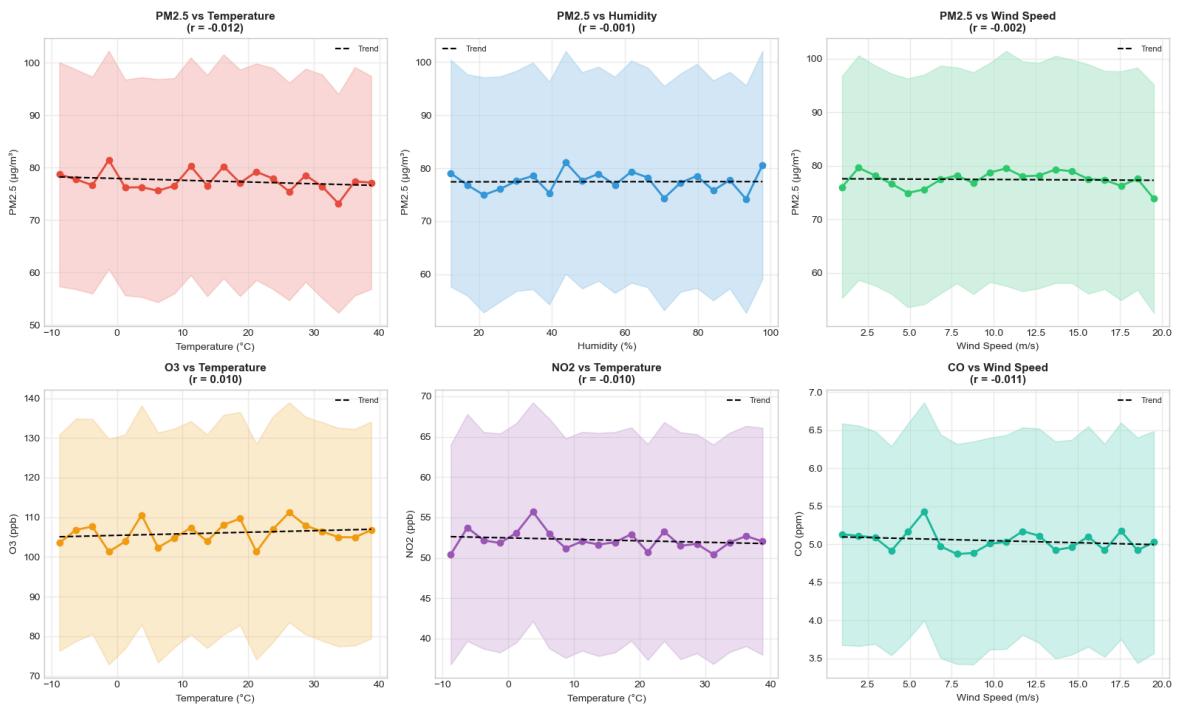
```

plt.suptitle('Bivariate Analysis - Pollutants vs Weather Conditions', fontsize=14)
plt.tight_layout()
plt.savefig('graphs/3_eda_bivariate/pollutants_vs_weather.png', dpi=150, bbox_inches='tight')
plt.show()

print("📊 Key Observations:")
print("-" * 50)
print(f" Temperature ↔ PM2.5: r = {df['Temperature'].corr(df['PM2.5']):.3f}")
print(f" Humidity ↔ PM2.5: r = {df['Humidity'].corr(df['PM2.5']):.3f}")
print(f" Wind Speed ↔ PM2.5: r = {df['Wind Speed'].corr(df['PM2.5']):.3f}")
print(f" Temperature ↔ O3: r = {df['Temperature'].corr(df['O3']):.3f}")
print(f" Temperature ↔ NO2: r = {df['Temperature'].corr(df['NO2']):.3f}")
print(f" Wind Speed ↔ CO: r = {df['Wind Speed'].corr(df['CO']):.3f}")

```

Bivariate Analysis - Pollutants vs Weather Conditions



📊 Key Observations:

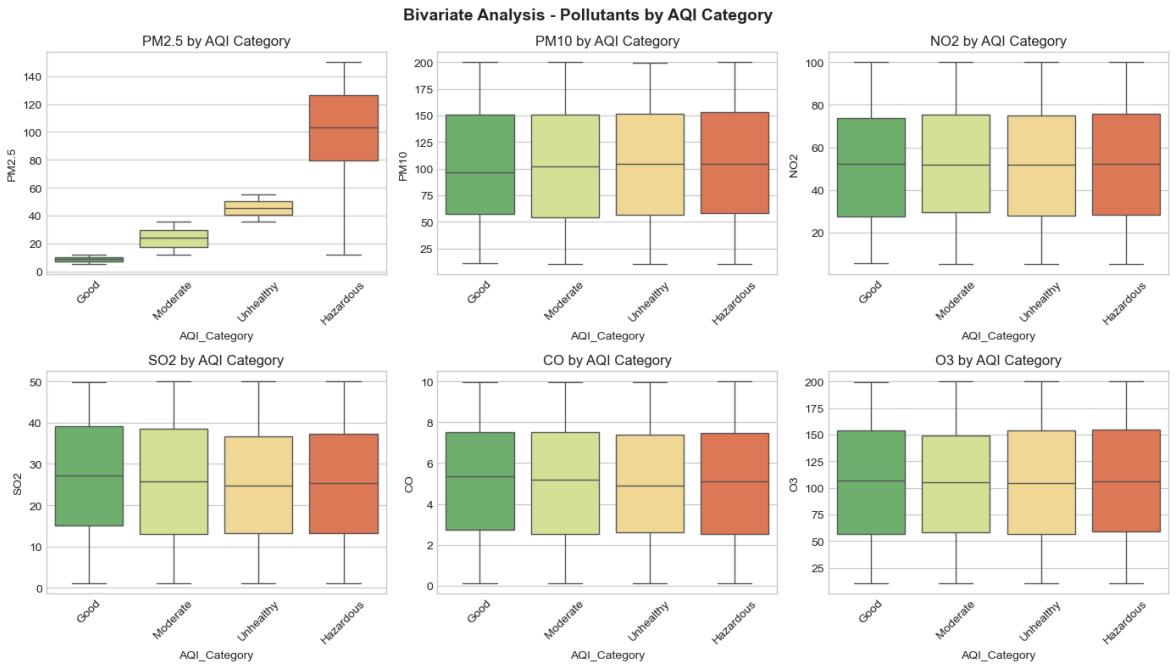
Temperature ↔ PM2.5: r = -0.012
 Humidity ↔ PM2.5: r = -0.001
 Wind Speed ↔ PM2.5: r = -0.002
 Temperature ↔ O3: r = 0.010
 Temperature ↔ NO2: r = -0.010
 Wind Speed ↔ CO: r = -0.011

```

In [19]: # 2.2b Bivariate - Pollutants by AQI Category
fig, axes = plt.subplots(2, 3, figsize=(14, 8))
axes = axes.flatten()
order = ['Good', 'Moderate', 'Unhealthy', 'Hazardous']

for i, col in enumerate(pollutant_cols):
    sns.boxplot(data=df, x='AQI_Category', y=col, ax=axes[i], order=order, palette=category10)
    axes[i].set_title(f'{col} by AQI Category')
    axes[i].tick_params(axis='x', rotation=45)
plt.suptitle('Bivariate Analysis - Pollutants by AQI Category', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/3_eda_bivariate/pollutants_by_aqi.png', dpi=150)
plt.show()

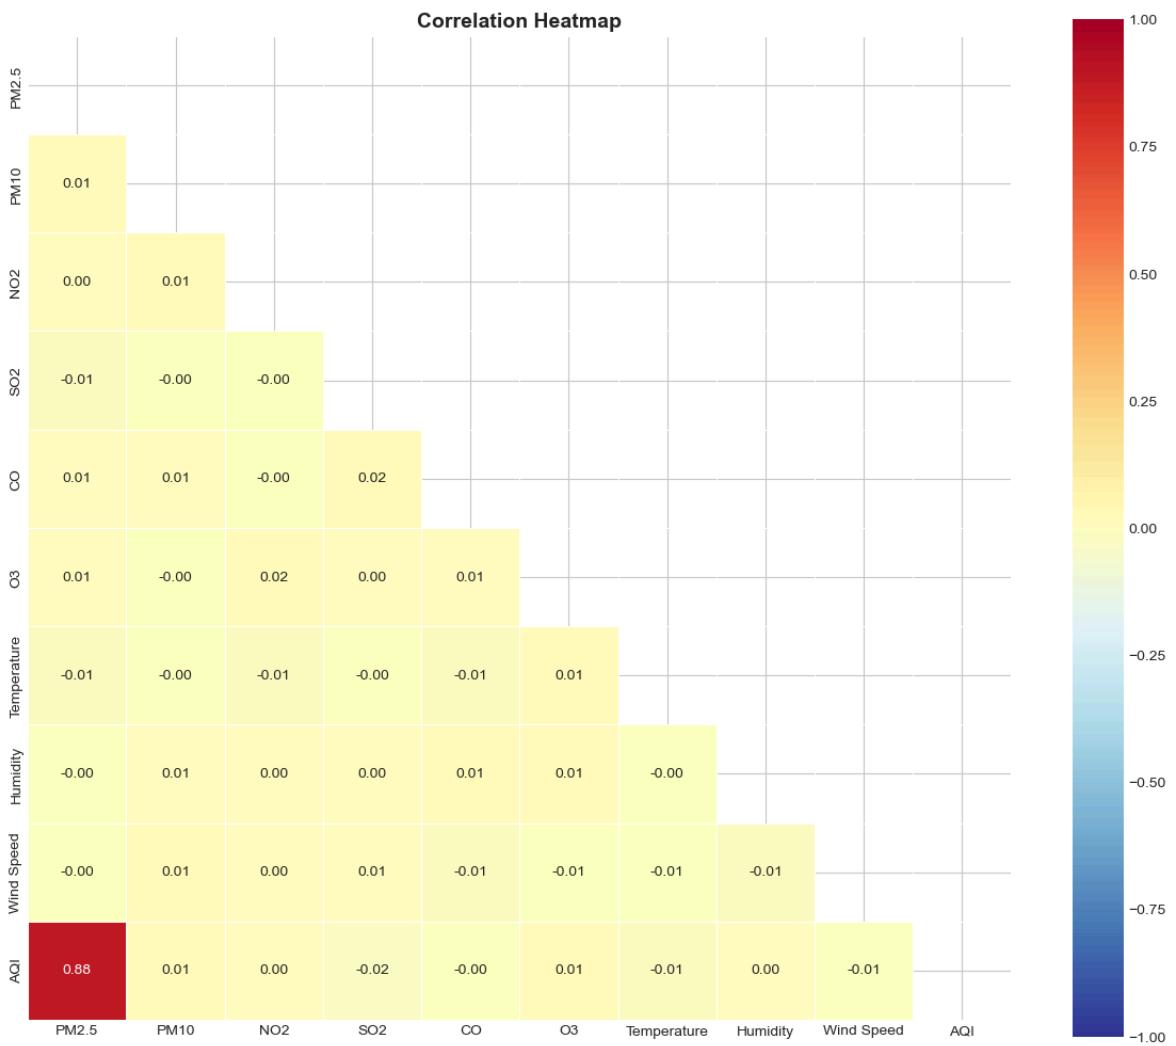
```



```
In [20]: # 2.3 Correlation Analysis
corr_cols = numerical_cols + ['AQI']
corr_matrix = df[corr_cols].corr()

plt.figure(figsize=(12, 10))
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
sns.heatmap(corr_matrix, mask=mask, annot=True, fmt='.2f', cmap='RdYlBu_r',
            center=0, square=True, linewidths=0.5, vmin=-1, vmax=1)
plt.title('Correlation Heatmap', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/4_correlation/correlation_heatmap.png', dpi=150)
plt.show()

# Strong correlations
print("Strong Correlations (|r| > 0.5):")
for i in range(len(corr_matrix.columns)):
    for j in range(i+1, len(corr_matrix.columns)):
        if abs(corr_matrix.iloc[i, j]) > 0.5:
            print(f" {corr_matrix.columns[i]} ↔ {corr_matrix.columns[j]}: {corr
```



Strong Correlations ($|r| > 0.5$):

PM2.5 \leftrightarrow AQI: 0.883

```
In [21]: # 2.4 Comparative Analysis - By City
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

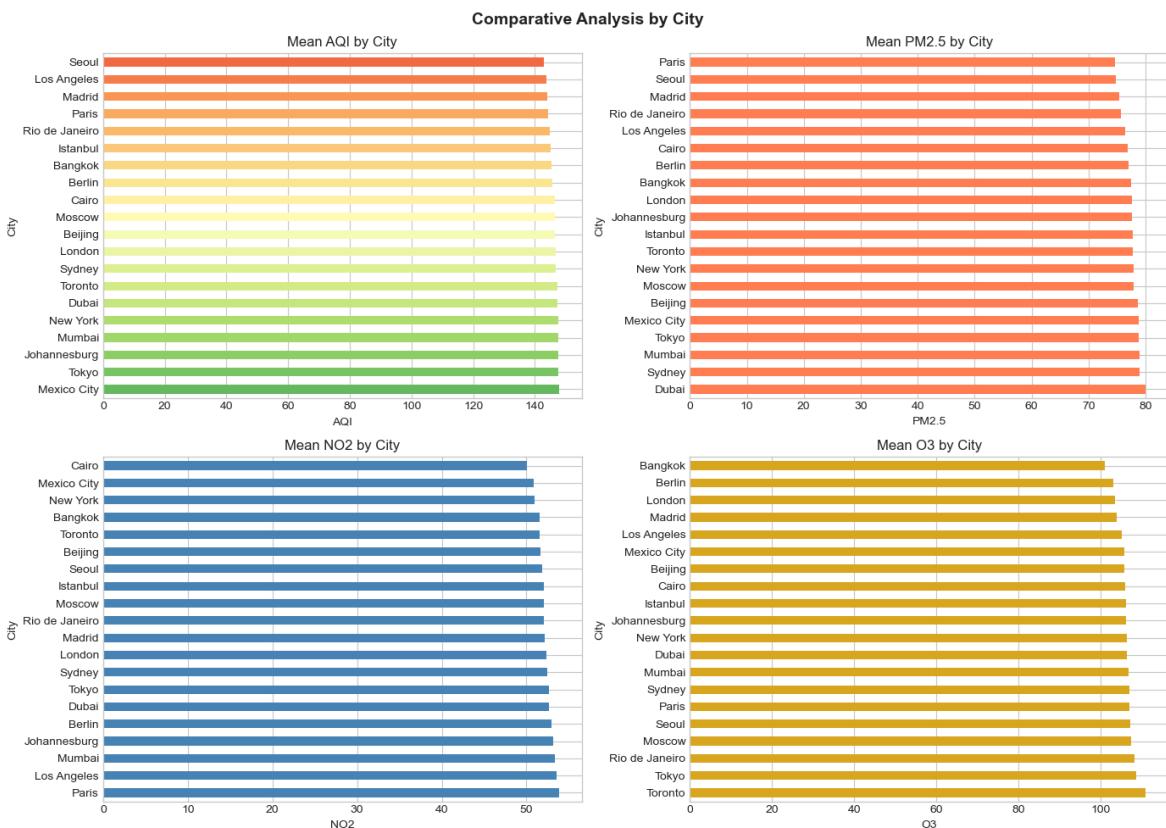
city_aqi = df.groupby('City')['AQI'].mean().sort_values(ascending=False)
city_aqi.plot(kind='barh', ax=axes[0,0], color=plt.cm.RdYlGn_r(np.linspace(0.2, 1, len(city_aqi))), axes[0,0].set_title('Mean AQI by City')); axes[0,0].set_xlabel('AQI')

city_pm25 = df.groupby('City')['PM2.5'].mean().sort_values(ascending=False)
city_pm25.plot(kind='barh', ax=axes[0,1], color='coral')
axes[0,1].set_title('Mean PM2.5 by City'); axes[0,1].set_xlabel('PM2.5')

city_no2 = df.groupby('City')['NO2'].mean().sort_values(ascending=False)
city_no2.plot(kind='barh', ax=axes[1,0], color='steelblue')
axes[1,0].set_title('Mean NO2 by City'); axes[1,0].set_xlabel('NO2')

city_o3 = df.groupby('City')['O3'].mean().sort_values(ascending=False)
city_o3.plot(kind='barh', ax=axes[1,1], color='goldenrod')
axes[1,1].set_title('Mean O3 by City'); axes[1,1].set_xlabel('O3')

plt.suptitle('Comparative Analysis by City', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/5_comparative/city_comparison.png', dpi=150)
plt.show()
```



```
In [22]: # 2.4b Comparative Analysis - By Country
```

```
print("Mean Pollutants by Country:")
print(df.groupby('Country')[pollutant_cols + ['AQI']].mean().round(2))
```

Mean Pollutants by Country:

Country	PM2.5	PM10	NO2	SO2	CO	O3	AQI
Australia	78.93	103.37	52.47	25.68	5.02	106.82	146.86
Brazil	75.67	105.44	52.10	24.58	4.91	108.07	144.86
Canada	77.83	103.87	51.55	24.80	4.95	110.85	147.27
China	78.63	103.19	51.62	24.90	5.09	105.74	146.48
Egypt	76.89	102.99	50.06	24.99	5.13	105.93	146.33
France	74.69	104.64	53.82	24.57	4.94	106.91	144.13
Germany	76.97	104.37	52.96	24.51	4.97	102.93	145.65
India	78.90	105.84	53.34	25.03	5.08	106.64	147.51
Japan	78.87	105.04	52.67	27.98	5.10	108.48	147.64
Mexico	78.86	104.53	50.86	26.02	5.14	105.69	147.73
Russia	77.88	105.94	52.03	26.11	5.12	107.29	146.34
South Africa	77.67	105.14	53.15	24.67	5.09	106.18	147.62
South Korea	74.80	105.00	51.88	25.01	4.89	107.02	142.83
Spain	75.43	105.87	52.14	25.15	4.89	103.78	143.89
Thailand	77.46	103.93	51.51	25.58	5.26	101.00	145.43
Turkey	77.71	103.14	52.02	25.77	5.12	106.18	145.16
UAE	80.01	103.89	52.68	24.96	5.08	106.40	147.40
UK	77.61	103.84	52.33	25.50	4.95	103.31	146.77
USA	77.11	104.29	52.30	25.67	5.12	105.69	145.50

```
In [23]: # 2.5 Time Series & Cycle Detection
```

```
df['Date'] = pd.to_datetime(df['Date'])
df['Year'] = df['Date'].dt.year
df['Month'] = df['Date'].dt.month
df['DayOfWeek'] = df['Date'].dt.dayofweek
df['Quarter'] = df['Date'].dt.quarter
```

```

month_names = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct']
day_names = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']

# Monthly Cycle
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

monthly_aqi = df.groupby('Month')['AQI'].mean()
axes[0,0].plot(monthly_aqi.index, monthly_aqi.values, marker='o', linewidth=2, color='blue')
axes[0,0].fill_between(monthly_aqi.index, monthly_aqi.values, alpha=0.3, color='blue')
axes[0,0].set_xticks(range(1, 13)); axes[0,0].set_xticklabels(month_names)
axes[0,0].set_title('Monthly AQI Cycle'); axes[0,0].set_ylabel('AQI')

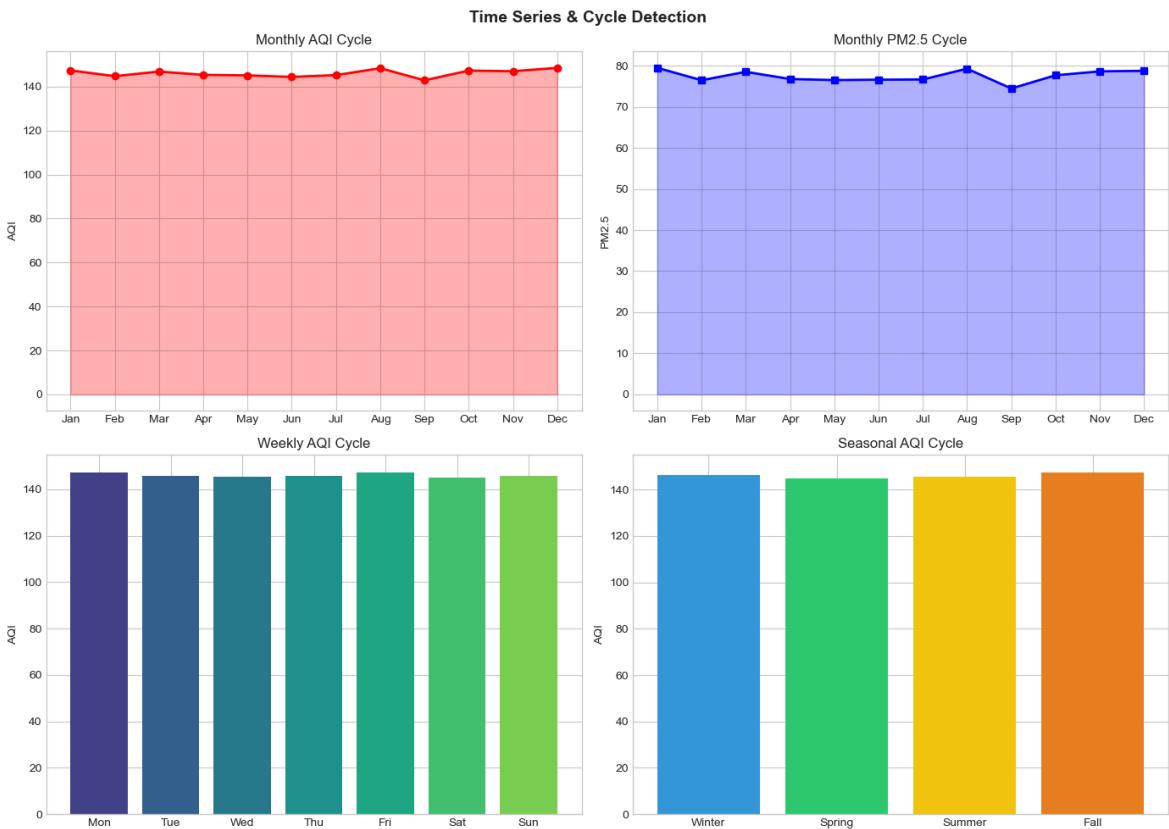
monthly_pm25 = df.groupby('Month')['PM2.5'].mean()
axes[0,1].plot(monthly_pm25.index, monthly_pm25.values, marker='s', linewidth=2, color='red')
axes[0,1].fill_between(monthly_pm25.index, monthly_pm25.values, alpha=0.3, color='red')
axes[0,1].set_xticks(range(1, 13)); axes[0,1].set_xticklabels(month_names)
axes[0,1].set_title('Monthly PM2.5 Cycle'); axes[0,1].set_ylabel('PM2.5')

# Weekly Cycle
weekly_aqi = df.groupby('DayOfWeek')['AQI'].mean()
axes[1,0].bar(weekly_aqi.index, weekly_aqi.values, color=plt.cm.viridis(np.linspace(0,1,7)))
axes[1,0].set_xticks(range(7)); axes[1,0].set_xticklabels(day_names)
axes[1,0].set_title('Weekly AQI Cycle'); axes[1,0].set_ylabel('AQI')

# Seasonal Cycle
season_map = {1: 'Winter', 2: 'Spring', 3: 'Summer', 4: 'Fall'}
df['Season'] = df['Quarter'].map(season_map)
seasonal_aqi = df.groupby('Season')['AQI'].mean().reindex(['Winter', 'Spring', 'Summer', 'Fall'])
axes[1,1].bar(seasonal_aqi.index, seasonal_aqi.values, color=['#3498db', '#2ecc71', '#f39c12', '#e74c3c'])
axes[1,1].set_title('Seasonal AQI Cycle'); axes[1,1].set_ylabel('AQI')

plt.suptitle('Time Series & Cycle Detection', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('graphs/6_timeseries/cycles.png', dpi=150)
plt.show()

```



3. Model Building & Prediction

```
In [24]: # Model Evaluation Function
label_map = {0: 'Good', 1: 'Moderate', 2: 'Unhealthy', 3: 'Hazardous'}
results = []

def train_evaluate_model(name, model):
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='weighted', zero_division=0)
    rec = recall_score(y_test, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)
    cv = cross_val_score(model, X_train_scaled, y_train, cv=5).mean()

    results.append({'Model': name, 'Accuracy': acc, 'Precision': prec, 'Recall': rec})

    print(f"\n{'*'*50}")
    print(f"MODEL: {name}")
    print(f"{'*'*50}")
    print(f"Accuracy: {acc:.4f} | Precision: {prec:.4f} | Recall: {rec:.4f} | F1: {f1:.4f}")
    print(f"\nConfusion Matrix:\n{confusion_matrix(y_test, y_pred)}")

return model, y_pred
```

```
In [36]: # Model 1: Logistic Regression
lr_model, lr_pred = train_evaluate_model('Logistic Regression', LogisticRegression())

fig, ax = plt.subplots(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, lr_pred), annot=True, fmt='d', cmap='Blues'
            xticklabels=label_map.values(), yticklabels=label_map.values())
```

```

plt.title('Logistic Regression - Confusion Matrix'); plt.xlabel('Predicted'); plt.tight_layout()
plt.savefig('graphs/7_models/logistic_regression_cm.png', dpi=150)
plt.show()

```

=====

MODEL: Logistic Regression

=====

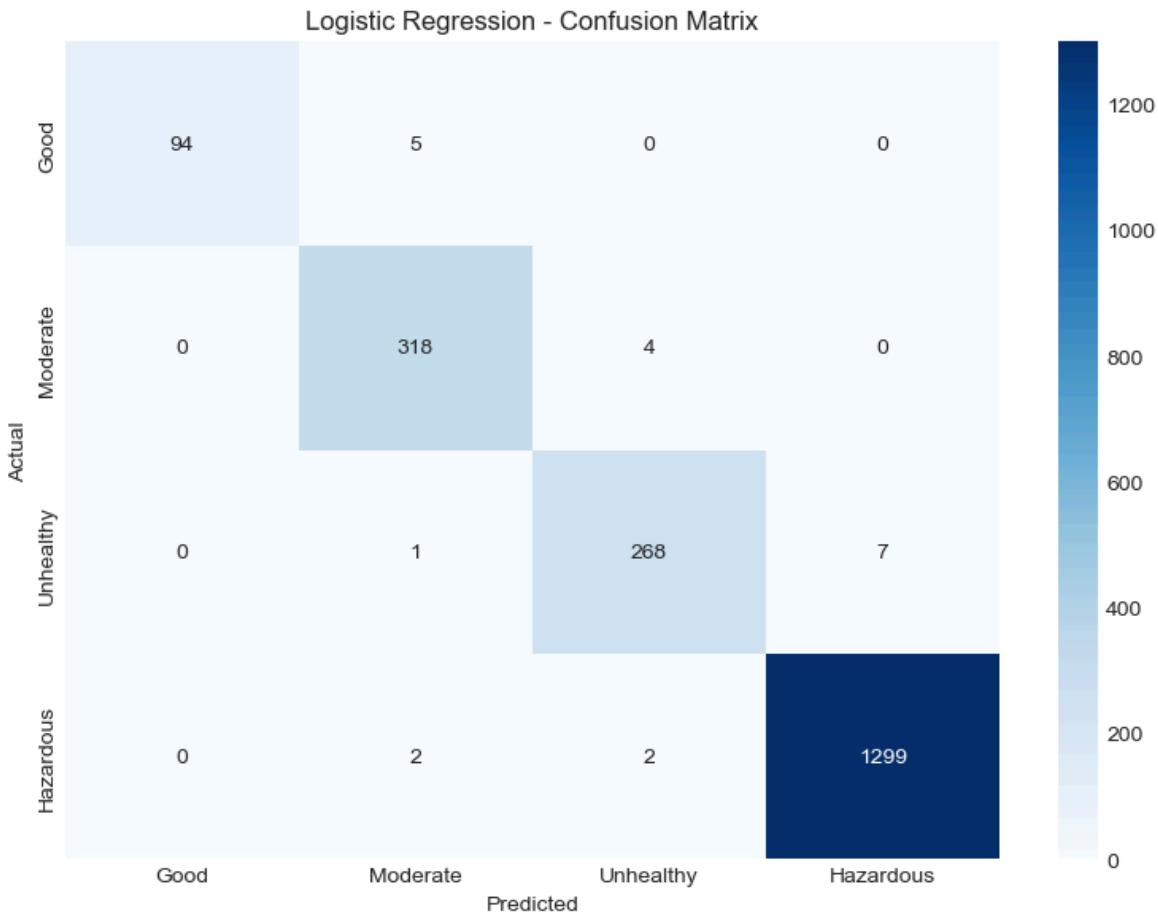
Accuracy: 0.9895 | Precision: 0.9895 | Recall: 0.9895 | F1: 0.9895

Confusion Matrix:

```

[[ 94   5   0   0]
 [  0 318   4   0]
 [  0   1 268   7]
 [  0   2   2 1299]]

```



In [37]:

```

# Model 2: Decision Tree
dt_model, dt_pred = train_evaluate_model('Decision Tree', DecisionTreeClassifier)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.heatmap(confusion_matrix(y_test, dt_pred), annot=True, fmt='d', cmap='Greens',
            xticklabels=label_map.values(), yticklabels=label_map.values())
axes[0].set_title('Decision Tree - Confusion Matrix'); axes[0].set_xlabel('Predicted')

dt_importance = pd.Series(dt_model.feature_importances_, index=features).sort_values()
dt_importance.plot(kind='barh', ax=axes[1], color='forestgreen')
axes[1].set_title('Decision Tree - Feature Importance')
plt.tight_layout()
plt.savefig('graphs/7_models/decision_tree.png', dpi=150)
plt.show()

```

```
=====
```

MODEL: Decision Tree

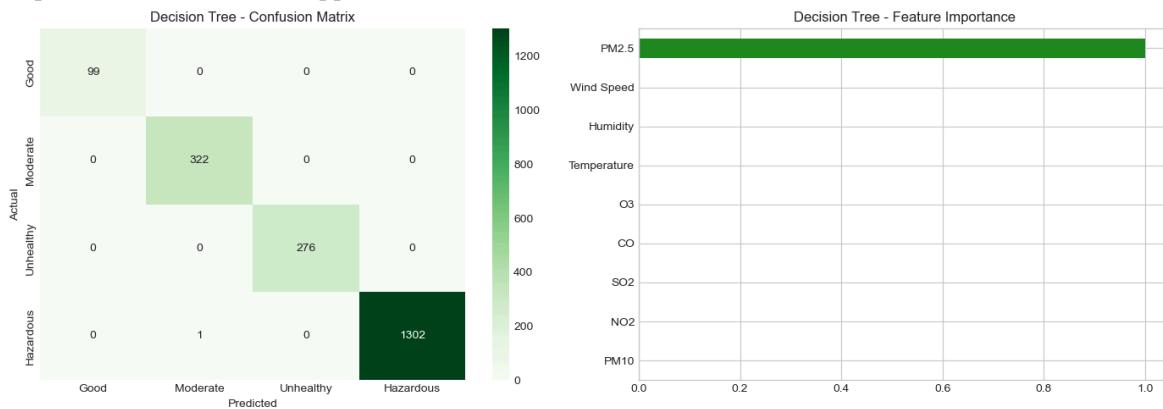
```
=====
```

Accuracy: 0.9995 | Precision: 0.9995 | Recall: 0.9995 | F1: 0.9995

```
=====
```

Confusion Matrix:

```
[[ 99   0   0   0]
 [  0 322   0   0]
 [  0   0 276   0]
 [  0   1   0 1302]]
```



```
In [46]: # Model 3: Random Forest
rf_model, rf_pred = train_evaluate_model('Random Forest', RandomForestClassifier)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
sns.heatmap(confusion_matrix(y_test, rf_pred), annot=True, fmt='d', cmap='YlGn',
            xticklabels=label_map.values(), yticklabels=label_map.values())
axes[0].set_title('Random Forest - Confusion Matrix'); axes[0].set_xlabel('Predicted')
axes[1].set_ylabel('Actual')

rf_importance = pd.Series(rf_model.feature_importances_, index=features).sort_values()
rf_importance.plot(kind='barh', ax=axes[1], color='darkgreen')
axes[1].set_title('Random Forest - Feature Importance')
plt.tight_layout()
plt.savefig('graphs/7_models/random_forest.png', dpi=150)
plt.show()
```

```
=====
```

MODEL: Random Forest

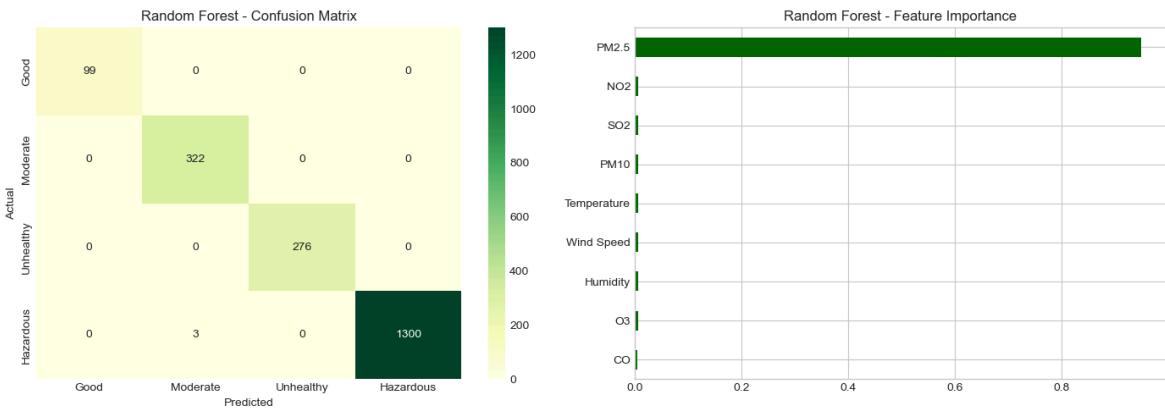
```
=====
```

Accuracy: 0.9985 | Precision: 0.9985 | Recall: 0.9985 | F1: 0.9985

```
=====
```

Confusion Matrix:

```
[[ 99   0   0   0]
 [  0 322   0   0]
 [  0   0 276   0]
 [  0   3   0 1300]]
```



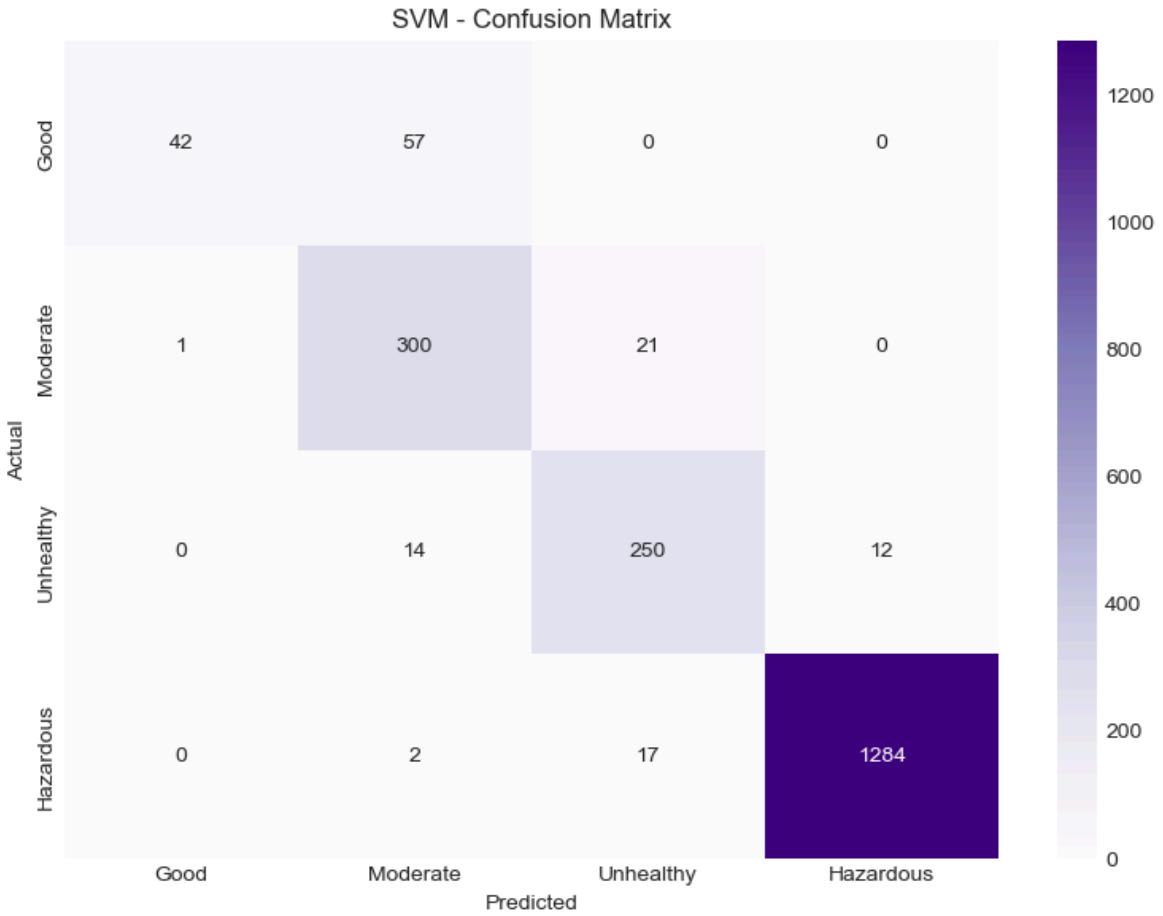
```
In [39]: # Model 4: Support Vector Machine (SVM)
svm_model, svm_pred = train_evaluate_model('SVM (RBF)', SVC(kernel='rbf', random_state=42), X_train, y_train)

fig, ax = plt.subplots(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, svm_pred), annot=True, fmt='d', cmap='Purples')
plt.title('SVM - Confusion Matrix'); plt.xlabel('Predicted'); plt.ylabel('Actual')
plt.tight_layout()
plt.savefig('graphs/7_models/svm_cm.png', dpi=150)
plt.show()

=====
MODEL: SVM (RBF)
=====
Accuracy: 0.9380 | Precision: 0.9431 | Recall: 0.9380 | F1: 0.9343
```

Confusion Matrix:

```
[[ 42  57   0   0]
 [  1 300  21   0]
 [  0  14 250  12]
 [  0    2  17 1284]]
```



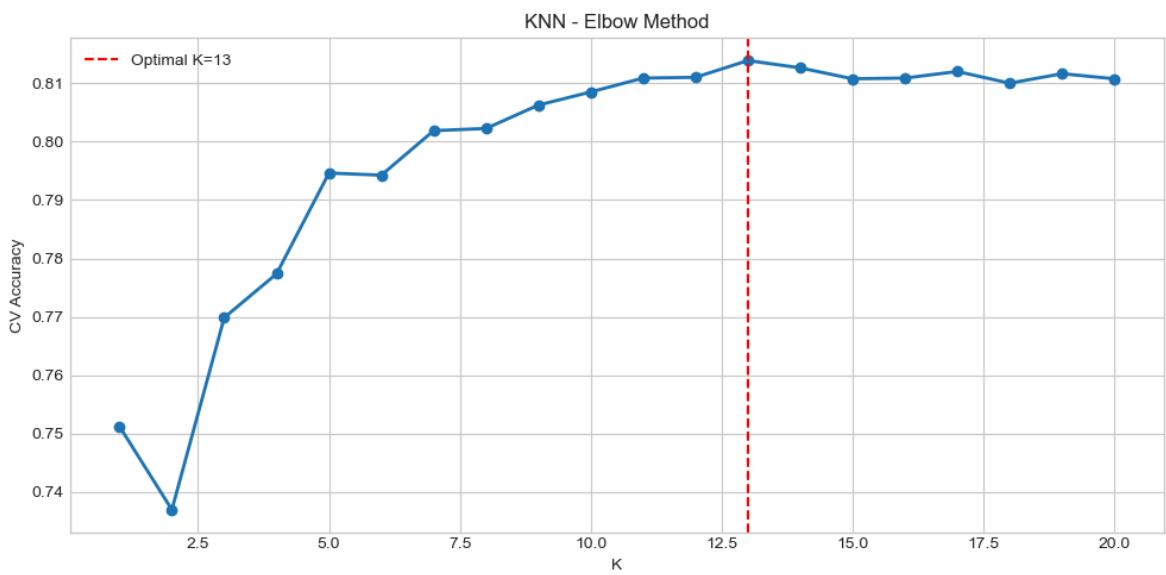
```
In [40]: # Model 5: K-Nearest Neighbors (KNN) with Elbow Method
k_scores = []
for k in range(1, 21):
    knn = KNeighborsClassifier(n_neighbors=k)
    cv_score = cross_val_score(knn, X_train_scaled, y_train, cv=5).mean()
    k_scores.append(cv_score)

optimal_k = np.argmax(k_scores) + 1

fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(range(1, 21), k_scores, marker='o', linewidth=2)
ax.axvline(optimal_k, color='red', linestyle='--', label=f'Optimal K={optimal_k}')
ax.set_xlabel('K'); ax.set_ylabel('CV Accuracy'); ax.set_title('KNN - Elbow Method')
ax.legend()
plt.tight_layout()
plt.savefig('graphs/7_models/knn_elbow.png', dpi=150)
plt.show()

knn_model, knn_pred = train_evaluate_model(f'KNN (K={optimal_k})', KNeighborsClassifer)

fig, ax = plt.subplots(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, knn_pred), annot=True, fmt='d', cmap='Oranges',
            xticklabels=label_map.values(), yticklabels=label_map.values())
plt.title('KNN - Confusion Matrix'); plt.xlabel('Predicted'); plt.ylabel('Actual')
plt.tight_layout()
plt.savefig('graphs/7_models/knn_cm.png', dpi=150)
plt.show()
```



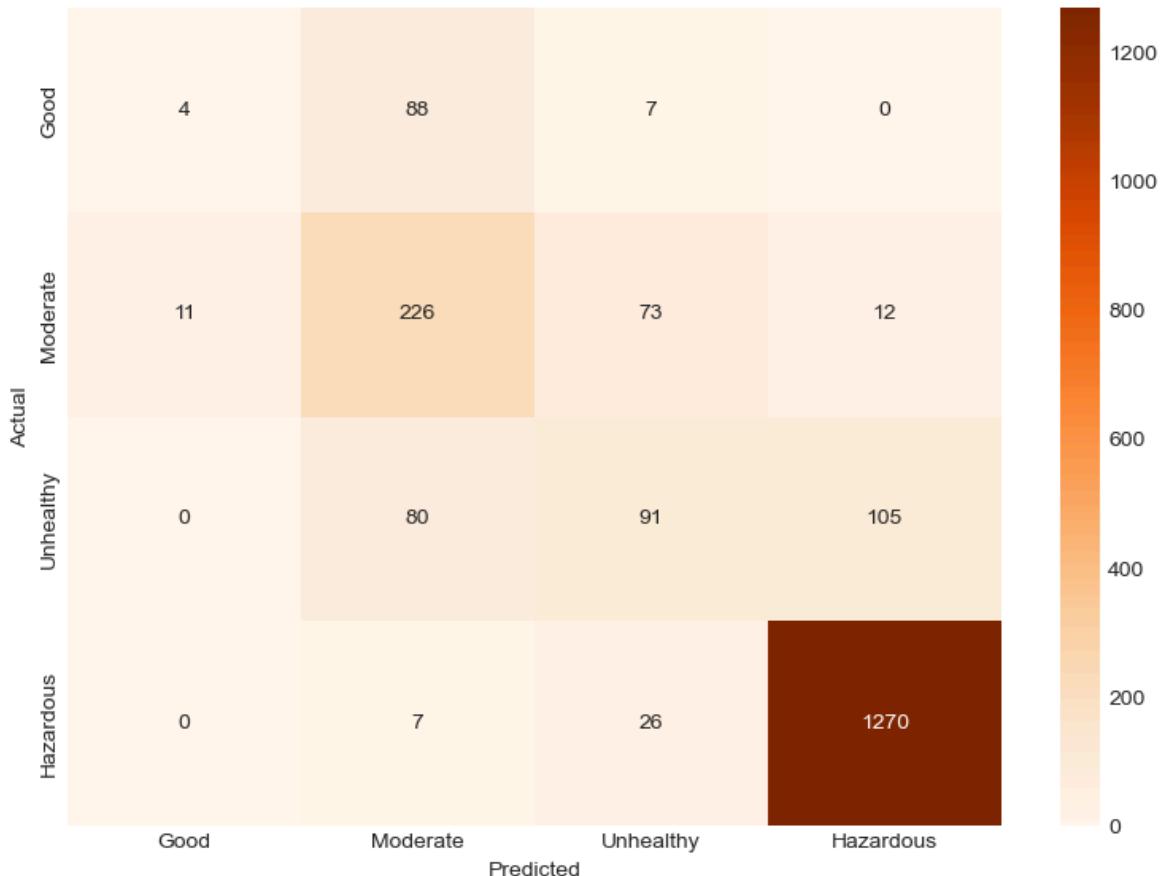
```
=====
MODEL: KNN (K=13)
=====
```

```
Accuracy: 0.7955 | Precision: 0.7642 | Recall: 0.7955 | F1: 0.7724
```

Confusion Matrix:

```
[[ 4   88    7    0]
 [ 11  226   73   12]
 [  0   80   91  105]
 [  0     7   26 1270]]
```

KNN - Confusion Matrix



```
In [41]: # Model Comparison
```

```
results_df = pd.DataFrame(results)
results_df = results_df.set_index('Model')
print("=" * 80)
```

```

print("MODEL COMPARISON SUMMARY")
print("=" * 80)
print(results_df.round(4).to_string())
print("\nBest Model:", results_df['F1'].idxmax(), f"(F1: {results_df['F1'].max()})")

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

results_df[['Accuracy', 'Precision', 'Recall', 'F1']].plot(kind='bar', ax=axes[0])
axes[0].set_title('Model Performance Metrics'); axes[0].set_ylabel('Score')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=45, ha='right')
axes[0].legend(loc='lower right'); axes[0].set_ylim(0, 1.1)

results_df['CV Score'].plot(kind='bar', ax=axes[1], color='steelblue', capsize=5)
axes[1].set_title('Cross-Validation Scores'); axes[1].set_ylabel('CV Accuracy')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45, ha='right')
axes[1].set_ylim(0, 1.1)

plt.tight_layout()
plt.savefig('graphs/7_models/model_comparison.png', dpi=150)
plt.show()

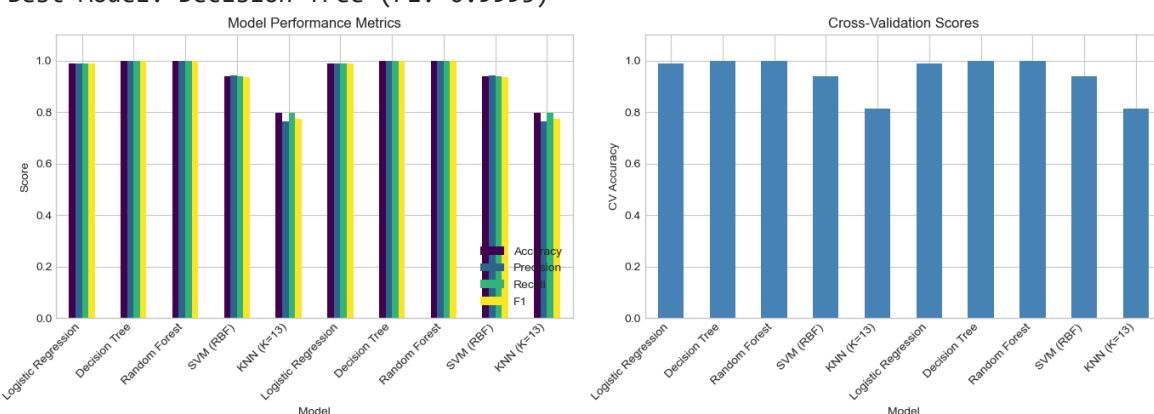
```

=====

MODEL COMPARISON SUMMARY

Model	Accuracy	Precision	Recall	F1	CV Score
Logistic Regression	0.9895	0.9895	0.9895	0.9895	0.9898
Decision Tree	0.9995	0.9995	0.9995	0.9995	0.9999
Random Forest	0.9985	0.9985	0.9985	0.9985	0.9991
SVM (RBF)	0.9380	0.9431	0.9380	0.9343	0.9391
KNN (K=13)	0.7955	0.7642	0.7955	0.7724	0.8139
Logistic Regression	0.9895	0.9895	0.9895	0.9895	0.9898
Decision Tree	0.9995	0.9995	0.9995	0.9995	0.9999
Random Forest	0.9985	0.9985	0.9985	0.9985	0.9991
SVM (RBF)	0.9380	0.9431	0.9380	0.9343	0.9391
KNN (K=13)	0.7955	0.7642	0.7955	0.7724	0.8139

Best Model: Decision Tree (F1: 0.9995)



```

In [ ]: # Feature Importance (from Random Forest)
feature_importance = pd.DataFrame({
    'Feature': features,
    'Importance': rf_model.feature_importances_
}).sort_values('Importance', ascending=True)

fig, ax = plt.subplots(figsize=(10, 6))
ax.barh(feature_importance['Feature'], feature_importance['Importance'], color='steelblue')
ax.set_xlabel('Importance'); ax.set_title('Random Forest - Feature Importance')

```

```

    for i, v in enumerate(feature_importance['Importance']):
        ax.text(v + 0.01, i, f'{v:.3f}', va='center')
plt.tight_layout()
plt.savefig('graphs/7_models/feature_importance.png', dpi=150)
plt.show()

print("\nTop 5 Most Important Features:")
print(feature_importance.tail().to_string(index=False))

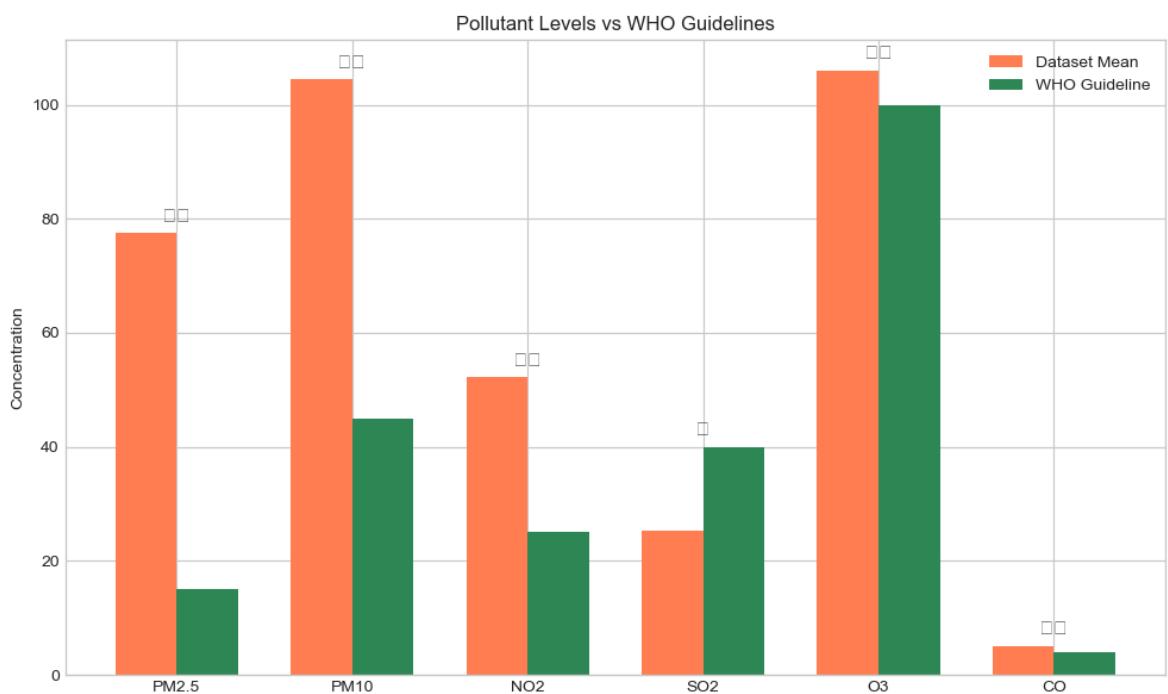
```

```

In [33]: # Health Impact Analysis - WHO Guidelines Comparison
who_guidelines = {'PM2.5': 15, 'PM10': 45, 'NO2': 25, 'SO2': 40, 'O3': 100, 'CO': 100}
actual_means = df[list(who_guidelines.keys())].mean()

fig, ax = plt.subplots(figsize=(10, 6))
x = np.arange(len(who_guidelines))
width = 0.35
ax.bar(x - width/2, actual_means.values, width, label='Dataset Mean', color='coral')
ax.bar(x + width/2, list(who_guidelines.values()), width, label='WHO Guideline',
       color='darkgreen')
ax.set_xticks(x); ax.set_xticklabels(who_guidelines.keys())
ax.set_ylabel('Concentration'); ax.set_title('Pollutant Levels vs WHO Guidelines')
ax.legend()
for i, (v1, v2) in enumerate(zip(actual_means.values, who_guidelines.values())):
    status = '⚠️' if v1 > v2 else '✓'
    ax.annotate(status, (i, max(v1, v2) + 2), ha='center', fontsize=12)
plt.tight_layout()
plt.savefig('graphs/5_comparative/who_comparison.png', dpi=150)
plt.show()

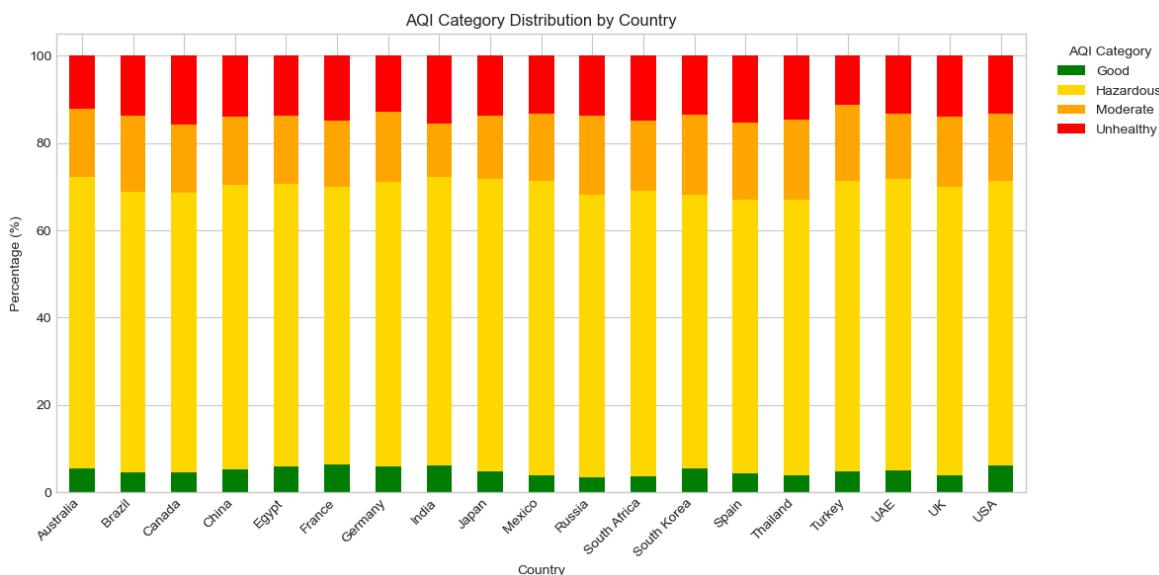
print("\nHealth Risk Assessment:")
for pollutant, who_val in who_guidelines.items():
    actual = actual_means[pollutant]
    pct = ((actual - who_val) / who_val) * 100
    status = "EXCEEDS" if actual > who_val else "WITHIN"
    print(f" {pollutant}: {actual:.2f} ({status} WHO limit by {abs(pct):.1f}%)"
```



Health Risk Assessment:

PM2.5: 77.45 (EXCEEDS WHO limit by 416.3%)
PM10: 104.44 (EXCEEDS WHO limit by 132.1%)
NO2: 52.20 (EXCEEDS WHO limit by 108.8%)
SO2: 25.34 (WITHIN WHO limit by 36.6%)
O3: 106.03 (EXCEEDS WHO limit by 6.0%)
CO: 5.05 (EXCEEDS WHO limit by 26.2%)

```
In [34]: # AQI Category Distribution by Region
fig, ax = plt.subplots(figsize=(12, 6))
aqi_by_country = df.groupby(['Country', 'AQI_Category']).size().unstack(fill_value=0)
aqi_by_country = aqi_by_country.div(aqi_by_country.sum(axis=1), axis=0) * 100
aqi_by_country.plot(kind='bar', stacked=True, ax=ax,
                     color=['green', 'gold', 'orange', 'red', 'purple', 'maroon'])
ax.set_ylabel('Percentage (%)'); ax.set_title('AQI Category Distribution by Country')
ax.legend(title='AQI Category', bbox_to_anchor=(1.02, 1))
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
plt.tight_layout()
plt.savefig('graphs/5_comparative/aqi_by_country.png', dpi=150)
plt.show()
```



```
In [ ]: # Summary Statistics and Key Findings
print("=" * 80)
print("PROJECT SUMMARY: GLOBAL AIR QUALITY ANALYSIS")
print("=" * 80)
print(f"\nDataset: {len(df)} records from {df['Country'].nunique()} countries, {df['Continent'].nunique()} continents")
print(f"\nAQI Distribution:")
for cat in df['AQI_Category'].value_counts().index:
    pct = (df['AQI_Category'] == cat).sum() / len(df) * 100
    print(f"  {cat}: {pct:.1f}%")

print(f"\nBest Performing Model: {results_df['F1-Score'].idxmax()}")
print(f"  - F1-Score: {results_df['F1-Score'].max():.4f}")
print(f"  - Accuracy: {results_df.loc[results_df['F1-Score'].idxmax(), 'Accuracy']}")

print(f"\nTop 3 Important Features:")
for _, row in feature_importance.tail(3).iterrows():
    print(f"  - {row['Feature']}: {row['Importance']:.4f}")

print("\nKey Recommendations:")
print("  1. Focus monitoring on PM2.5 and PM10 as primary health risk indicators")
print("  2. Implement targeted interventions in regions with 'Unhealthy' or worse AQI categories")
```

```
print(" 3. Use ML models for real-time AQI prediction and early warning systems")
print(" 4. Weather conditions (temp, humidity) significantly influence pollutant")
print("=" * 80)
```

```
In [ ]: # Save Final Results to CSV
results_df.to_csv('model_results.csv')
feature_importance.to_csv('feature_importance.csv', index=False)
print("Results saved to CSV files!")
print(f"\nAll graphs saved to organized folders:")
for folder in sorted(os.listdir('graphs')):
    files = os.listdir(f'graphs/{folder}')
    print(f"  └── graphs/{folder}/ ({len(files)}) files")
    for f in files:
        print(f"    └── {f}")
```